

Hadoop 2.2.0 installation

Chun-Chen Tu
timtu@umich.edu

Before installation

- Where to get hadoop 2.2.0
 - <http://apache.stu.edu.tw/hadoop/common/hadoop-2.2.0/>
 - ftp://hadoop:hahahadoop@140.113.114.104/hadoop_2.2.0_cluster.tar.gz
- GUI mode may help for typing commands.
- In this ppt, commands will be shown in italic and purple color
 - *mkdir hadoop*
- The content in file will be label as green in a square
 - Hello, Hadoop

About Hadoop

- If you download from official website, you should check the library files

cd hadoop/lib/native

*file **

```
hadoop@ubuntu:~/Downloads/tmp/hadoop-2.2.0/lib/native$ file *.so.*
libhadoop.so.1.0.0: ELF 32-bit LSB shared object, Intel 80386, version 1 (SYSV)
not stripped, BuildID[sha1]=0x9e1d49b05f67d38454e42b216e053a27ae8bac9, not stripped
libhdfs.so.0.0.0: ELF 32-bit LSB shared object, Intel 80386, version 1 (SYSV)
not stripped, BuildID[sha1]=0x5f6c9ec21598be534211b48fc692e7858f21afc, not stripped
```


- The default library is designed for 32-bit system. There will be some warnings but Hadoop can still work. To fix it, re-compile Hadoop or get a 64-bit library.

Brief introduction

- Goal : setting up cluster of Hadoop 2.2.0
- Workflow:
 - Download and install required material.
 - Install JAVA
 - Install Hadoop
 - Cluster configuration
 - Start and test

Topology setting

- 1 machine running namenode
- 1 machine running secondary namenode
- 1 resource manager
- 1 client running job history server
- other 16 machines running datanode and node manager.



NOTE: You can make up host yourselves.

| Hostname | IP | Role |
|----------|---------------|-------------------------------|
| cloud01 | 192.168.1.101 | Namenode |
| cloud02 | 192.168.1.102 | Secondary Namenode |
| cloud03 | 192.168.1.103 | Resource manager |
| cloud04 | 192.168.1.104 | Client, Job history server |
| cloud05 | 192.168.1.105 | Datanode, Nodemanager |
| cloud20 | 192.168.1.120 | |

Notes

- We create a user called “hadoop” and we want to install hadoop under this user.
- Some settings should be exactly same for all machines. We will use shell scripts to make things easy.
- First, we will set up Namenode and copy the settings to other machines.
- Let's start with cloud01 !!

After login, install required packages first

sudo apt-get install libssl-dev rsync g++

type “y” when asked

Note: If you get message like “Package not found” type

sudo apt-get update

```
hadoop@ubuntu:~$ ls
Desktop  Documents  Downloads  Music  Pictures  Public  Templates  Videos
hadoop@ubuntu:~$ sudo apt-get install libssl-dev rsync g++
[sudo] password for hadoop:
Reading package lists... Done
Building dependency tree
Reading state information... Done
rsync is already the newest version.
The following extra packages will be installed:
  g++-4.6 libssl-doc libstdc++6-4.6-dev zlib1g-dev
Suggested packages:
  g++-multilib g++-4.6-multilib gcc-4.6-doc libstdc++6-4.6-dbg
  libstdc++6-4.6-doc
The following NEW packages will be installed:
  g++ g++-4.6 libssl-dev libssl-doc libstdc++6-4.6-dev zlib1g-dev
0 upgraded, 6 newly installed, 0 to remove and 61 not upgraded.
Need to get 11.4 MB of archives.
After this operation, 33.5 MB of additional disk space will be used.
Do you want to continue [Y/n]?
```


Edit hosts files

sudo vim /etc/hosts

Add the hostname and ip for each machines in the cluster

DELETE two lines of 127.x.x.x

Thus, it will look like this

| | |
|---------------|---------|
| 192.168.1.101 | cloud01 |
| 192.168.1.102 | cloud02 |
| 192.168.1.103 | cloud03 |
| 192.168.1.104 | cloud04 |
| 192.168.1.105 | cloud05 |
| 192.168.1.106 | cloud06 |
| 192.168.1.107 | cloud07 |
| 192.168.1.108 | cloud08 |
| 192.168.1.109 | cloud09 |
| 192.168.1.110 | cloud10 |
| 192.168.1.111 | cloud11 |
| 192.168.1.112 | cloud12 |
| 192.168.1.113 | cloud13 |
| 192.168.1.114 | cloud14 |
| 192.168.1.115 | cloud15 |
| 192.168.1.116 | cloud16 |
| 192.168.1.117 | cloud17 |
| 192.168.1.118 | cloud18 |
| 192.168.1.119 | cloud19 |
| 192.168.1.120 | cloud20 |

You should add this information on **all** machines.

Download files:

cd ~/

wget ftp://hadoop:hahahadoop@140.113.114.104/hadoop_2.2.0_cluster.tar.gz

wget ftp://hadoop:hahahadoop@140.113.114.104/jdk-7u45-linux-x64.gz

```
Connecting to 140.113.114.104:21... connected.
Logging in as hadoop ... Logged in!
==> SYST ... done.      ==> PWD ... done.
==> TYPE I ... done.    ==> CWD not needed.
==> SIZE hadoop-1.2.1.tar.gz ... 63851630
==> PASV ... done.      ==> RETR hadoop-1.2.1.tar.gz ... done.
Length: 63851630 (61M) (unauthoritative)

100%[=====>] 63,851,630  81.2M/s   in 0.8s

2013-12-29 12:32:57 (81.2 MB/s) - `hadoop-1.2.1.tar.gz' saved [63851630]

hadoop@ubuntu:~/Downloads$ wget ftp://hadoop:hahahadoop@140.113.114.104/jdk-7u45-
linux-x64.gz
--2013-12-29 12:34:05--  ftp://hadoop:*password*@140.113.114.104/jdk-7u45-linux-
x64.gz
=> `jdk-7u45-linux-x64.gz'
Connecting to 140.113.114.104:21... connected.
Logging in as hadoop ... Logged in!
==> SYST ... done.      ==> PWD ... done.
==> TYPE I ... done.    ==> CWD not needed.
==> SIZE jdk-7u45-linux-x64.gz ... 138094686
==> PASV ... done.      ==> RETR jdk-7u45-linux-x64.gz ... done.
Length: 138094686 (132M) (unauthoritative)

100%[=====>] 138,094,686  85.3M/s   in 1.5s

2013-12-29 12:34:07 (85.3 MB/s) - `jdk-7u45-linux-x64.gz' saved [138094686]

hadoop@ubuntu:~/Downloads$
```

Install java : reference [website](#)

(Under Downloads folder)

```
tar -zxvf jdk-7u45-linux-x64.gz
```


```
sudo mkdir /usr/lib/jdk
```

```
sudo cp -r jdk1.7.0_45 /usr/lib/jdk/
```

Edit profile:

```
sudo vim ~/.bashrc
```

(add four lines in at the top of .bashrc)



```
export JAVA_HOME=/usr/lib/jdk/jdk1.7.0_45
export JRE_HOME=/usr/lib/jdk/jdk1.7.0_45/jre
export PATH=$JAVA_HOME/bin:$JAVA_HOME/jre/bin:$PATH
export CLASSPATH=$CLASSPATH:.$JAVA_HOME/lib:$JAVA_HOME/jre/lib
```

```
export JAVA_HOME=/usr/lib/jdk/jdk1.7.0_45
export JRE_HOME=/usr/lib/jdk/jdk1.7.0_45/jre
export PATH=$JAVA_HOME/bin:$JAVA_HOME/jre/bin:$PATH
export CLASSPATH=$CLASSPATH:.$JAVA_HOME/lib:$JAVA_HOME/jre/lib
```

```
source ~/.bashrc
```

Config java:

```
sudo update-alternatives --install /usr/bin/java java /usr/lib/jdk/jdk1.7.0_45/bin/java 300
```

```
sudo update-alternatives --install /usr/bin/javac javac /usr/lib/jdk/jdk1.7.0_45/bin/javac 300
```

```
sudo update-alternatives --config java
```

```
sudo update-alternatives --config javac
```

Test it with version

```
java -version
```

You will see the version information if success.

```
hadoop@ubuntu:~/Downloads$ java -version
java version "1.7.0_45"
Java(TM) SE Runtime Environment (build 1.7.0_45-b18)
Java HotSpot(TM) 64-Bit Server VM (build 24.45-b08, mixed mode)
hadoop@ubuntu:~/Downloads$
```

SSH setting: SSH setting is optional but is recommended if you don't want to enter password every time.

Generate RSA key

```
ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
```

put public key on current machine

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

Copy public key to other machines

```
ssh hadoop@cloud02 "mkdir ~/.ssh"
```

```
scp ~/.ssh/id_rsa.pub hadoop@cloud02:~/.ssh/keys_from_hosts
```

```
ssh hadoop@cloud02 "cat ~/.ssh/keys_from_hosts >> ~/.ssh/authorized_keys"
```

```
hadoop@ubuntu:~/Downloads$ ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
Generating public/private rsa key pair.
Created directory '/home/hadoop/.ssh'.
Your identification has been saved in /home/hadoop/.ssh/id_rsa.
Your public key has been saved in /home/hadoop/.ssh/id_rsa.pub.
The key fingerprint is:
eb:28:46:7d:29:87:08:63:cd:d0:f5:a8:76:82:13:2a hadoop@ubuntu
The key's randomart image is:
+--[ RSA 2048 ]-----+
|      . . .      |
|      . . 0      |
|      .+ . .      |
|     .+00.        |
|E.00+0..S.        |
|    000+ +.       |
|      . +.        |
|      0 0         |
|      . . . .     |
+-----+

```

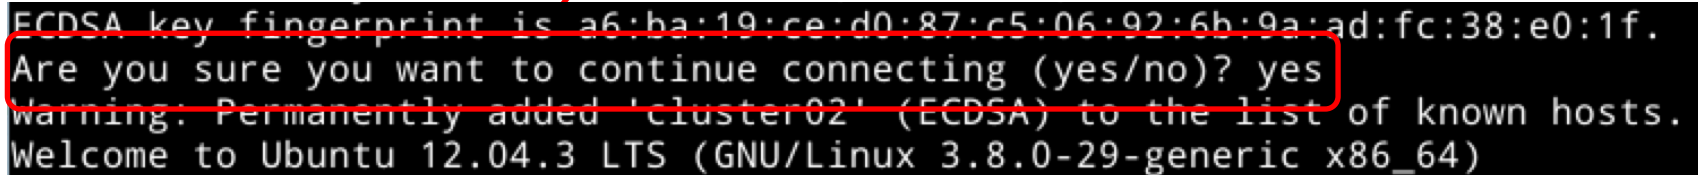
SSH test:

ssh hadoop@cloud02

remember to exit

exit

You will be asked for the authenticity for the first time. After this connection, no more inquiring.

A terminal window showing the SSH connection process. The text is as follows:
ECDSA key fingerprint is a6:ba:19:ce:d0:87:c5:06:92:6b:9a:ad:fc:38:e0:1f.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'cluster02' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 12.04.3 LTS (GNU/Linux 3.8.0-29-generic x86_64)
A red box highlights the prompt "Are you sure you want to continue connecting (yes/no)? yes" and the response "yes". A red arrow points from the text "You will be asked for the authenticity for the first time. After this connection, no more inquiring." to the red box.

```
ECDSA key fingerprint is a6:ba:19:ce:d0:87:c5:06:92:6b:9a:ad:fc:38:e0:1f.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'cluster02' (ECDSA) to the list of known hosts.  
Welcome to Ubuntu 12.04.3 LTS (GNU/Linux 3.8.0-29-generic x86_64)
```

```
* Documentation:  https://help.ubuntu.com/
```

```
System information as of Sun Jan  5 14:59:37 CST 2014
```

| | | | |
|---------------|------------------|----------------------|-----------|
| System load: | 0.01 | Processes: | 127 |
| Usage of /: | 88.6% of 12.71GB | Users logged in: | 1 |
| Memory usage: | 28% | IP address for eth0: | 10.0.2.15 |
| Swap usage: | 0% | | |

```
=> / is using 88.6% of 12.71GB
```

```
=> There is 1 zombie process.
```

```
Graph this data and manage this system at https://landscape.canonical.com/
```

```
67 packages can be updated.
```

```
32 updates are security updates.
```

If you fail the setting, you will need to enter password.

Set up ssh of all the machines in cluster using shell scripts

cd ~/hadoop/scripts (This directory will appear after tar hadoop_1.2.1_cluster.tar.gz

1. List of machines See next slide)

vim machines

```
cloud02
cloud03
cloud04
cloud05
cloud06
cloud07
cloud08
cloud09
cloud10
cloud11
cloud12
cloud13
cloud14
cloud15
cloud16
cloud17
cloud18
cloud19
cloud20
```

2. And then we create a shell script SetSSH.sh to do jobs according to the list of machines

vim SetSSH.sh

```
#!/bin/bash
HOST_FILES=/home/hadoop/hadoop/scriptes/machines
seq=1
while read line
do
    lines[$seq]=$line
    ((seq++))
done < $HOST_FILES

for ((i=1;i<=${#lines[@]};i++))
do
    echo "Set keys to ${lines[i]}"
    ssh ${lines[i]} "mkdir ~/.ssh"
    scp -r ~/.ssh/id_rsa.pub ${lines[i]}:~/.ssh/keys_from_hosts
    ssh ${lines[i]} "cat ~/.ssh/keys_from_hosts >> ~/.ssh/authorized_keys"
done
```

3. Finally change scripts to executable and run it

chmod 755 setSSH.sh

./setSSH.sh

Install hadoop:

tar -zxvf hadoop_2.2.0_cluster.tar.gz

mv hadoop ~/hadoop move it under home directory for convenience

vim ~/.bashrc

```
export HADOOP_HOME=/home/hadoop/hadoop
export HADOOP_PREFIX=/home/hadoop/hadoop
export HADOOP_COMMON_HOME=/home/hadoop/hadoop
export HADOOP_MAPRED_HOME=/home/hadoop/hadoop
export HADOOP_CONF_DIR=/home/hadoop/hadoop/etc/hadoop
export HADOOP_HDFS_HOME=/home/hadoop/hadoop
export HADOOP_YARN_HOME=/home/hadoop/hadoop
export YARN_CONF_DIR=/home/hadoop/hadoop/etc/hadoop

export HADOOP_COMMON_LIB_NATIVE_DIR=/home/hadoop/hadoop/lib/native
export HADOOP_OPTS="-Djava.library.path=/home/hadoop/hadoop/lib/native"
```

source ~/.bashrc

Configure for cluster: you should set up 4 files in `hadoop/etc/hadoop/`
`core-site.xml` `hdfs-site.xml` `mapred-site.xml` `yarn-site.xml`

`core-site.xml` : parameter [website](#)

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://cloud01:9000</value>
  </property>
  <property>
    <name>io.file.buffer.size</name>
    <value>131072</value>
  </property>
  <property>
    <name>hadoop.tmp.dir</name>
    <value>/home/hadoop/tmp</value>
  </property>
  <property>
    <name>hadoop.proxyuser.hadoop.hosts</name>
    <value>*</value>
  </property>
  <property>
    <name>hadoop.proxyuser.hadoop.groups</name>
    <value>*</value>
  </property>
</configuration>
```


hdfs-site.xml : parameter [website](#)

```
<configuration>
  <property>
    <name>dfs.namenode.secondary.http-address</name>
    <value>cloud02:9001</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>/home/hadoop/dfs/name</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>/home/hadoop/dfs/data</value>
  </property>
  <property>
    <name>dfs.replication</name>
    <value>3</value>
  </property>
  <property>
    <name>dfs.webhdfs.enabled</name>
    <value>true</value>
  </property>
</configuration>
```

mapred-site.xml : parameter [website](#)

```
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
  <property>
    <name>mapreduce.jobhistory.address</name>
    <value>cloud04:10020</value>
  </property>
  <property>
    <name>mapreduce.jobhistory.webapp.address</name>
    <value>cloud04:19888</value>
  </property>
</configuration>
```

yarn-site.xml : parameter [website](#)

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>
  </property>
  <property>
    <name>yarn.resourcemanager.address</name>
    <value>cloud03:8032</value>
  </property>
  <property>
    <name>yarn.resourcemanager.scheduler.address</name>
    <value>cloud03:8030</value>
  </property>
  <property>
    <name>yarn.resourcemanager.resource-tracker.address</name>
    <value>cloud03:8031</value>
  </property>
  <property>
    <name>yarn.resourcemanager.admin.address</name>
    <value>cloud03:8033</value>
  </property>
  <property>
    <name>yarn.resourcemanager.webapp.address</name>
    <value>cloud03:8088</value>
  </property>
</configuration>
```

Similarly under hadoop/etc/hadoop/
edit slaves and masters

In general, masters contain namenode, secondary namenode, resourcemanager and job history server.

In our case, in masters we add

```
cloud01  
cloud02  
cloud03  
cloud04
```

slaves will include datanode and nodemanager
in slaves we add

```
cloud05  
cloud06  
:  
cloud20
```

Tar the current hadoop directory and copy files to other machine using shell scripts cpHadoop.sh

cd ~/

rm hadoop_2.2.0_cluster.tar.gz

tar -czvf hadoop_2.2.0_cluster.tar.gz hadoop

cd ~/hadoop/scripts/

```
dir=/home/hadoop/hadoop/etc/hadoop
```

```
HOST_FILES=/home/hadoop/hadoop/scripts/machines
```

```
seq=1
```

```
while read line
```

```
do
```

```
    lines[$seq]=$line
```

```
    ((seq++))
```

```
done < $HOST_FILES
```

```
for ((i=1;i<=${#lines[@]};i++))
```

```
do
```

```
    scp ~/Downloads/hadoop_2.2.0_cluster.tar.gz ${lines[$i]}:~/
```

```
    ssh ${lines[$i]} "tar -zxf hadoop_2.2.0_cluster.tar.gz"
```

```
    echo "Copy file: .bashrc to: ${lines[$i]}"
```

```
    scp ~/.bashrc ${lines[$i]}:~/.bashrc
```

```
for FILE in slaves masters core-site.xml hdfs-site.xml mapred-site.xml yarn-site.xml
```

```
do
```

```
    echo "Copy file: $FILE to: ${lines[$i]}"
```

```
    scp $dir/$FILE ${lines[$i]}:$dir/$FILE
```

```
done
```

```
done
```

HDFS format:

hdfs namenode -format

```
STARTUP_MSG:   java = 1.7.0_45
*****/
13/12/29 12:50:56 INFO util.GSet: Computing capacity for map BlocksMap
13/12/29 12:50:56 INFO util.GSet: VM type           = 64-bit
13/12/29 12:50:56 INFO util.GSet: 2.0% max memory = 1013645312
13/12/29 12:50:56 INFO util.GSet: capacity         = 2^21 = 2097152 entries
13/12/29 12:50:56 INFO util.GSet: recommended=2097152, actual=2097152
13/12/29 12:50:57 INFO namenode.FSNamesystem: fsOwner=hadoop
13/12/29 12:50:57 INFO namenode.FSNamesystem: supergroup=supergroup
13/12/29 12:50:57 INFO namenode.FSNamesystem: isPermissionEnabled=true
13/12/29 12:50:57 INFO namenode.FSNamesystem: dfs.block.invalidate.limit=100
13/12/29 12:50:57 INFO namenode.FSNamesystem: isAccessTokenEnabled=false accessK
eyUpdateInterval=0 min(s), accessTokenLifetime=0 min(s)
13/12/29 12:50:57 INFO namenode.FSEditLog: dfs.namenode.edits.toleration.length
= 0
13/12/29 12:50:57 INFO namenode.NameNode: Caching file names occurring more than
10 times
13/12/29 12:50:57 INFO common.Storage: Image file /tmp/hadoop-hadoop/dfs/name/cu
rrent/fsimage of size 112 bytes saved in 0 seconds.
13/12/29 12:50:57 INFO namenode.FSEditLog: closing edit log: position=4, editlog
=/tmp/hadoop-hadoop/dfs/name/current/edits
13/12/29 12:50:57 INFO namenode.FSEditLog: close success: truncate to 4, editlog
=/tmp/hadoop-hadoop/dfs/name/current/edits
13/12/29 12:50:57 INFO common.Storage: Storage directory /tmp/hadoop-hadoop/dfs/
name has been successfully formatted.
13/12/29 12:50:57 INFO namenode.NameNode: SHUTDOWN_MSG:
*****/
SHUTDOWN_MSG: Shutting down NameNode at ubuntu/127.0.1.1
*****/
```

Start hadoop
cd ~/hadoop/sbin
./start-all.sh

jps : see what's working on current machine.
hdfs dfsadmin -report : see the information of DFS.
more commands on this [website](#)

Start resource manager
ssh cloud03 "~/hadoop/sbin/yarn-start.sh start resourcemanager"
check if it start
ssh cloud03 "jps"

Start job history server
ssh cloud04 "~/hadoop/sbin/mr-jobhistory-daemon.sh start historyserver"
check if it start
ssh cloud04 "jps"

Let's run an example! There is an example jar file under
hadoop/share/hadoop/mapreduce : hadoop-mapreduce-examples-2.2.0.jar

hadoop jar hadoop-mapreduce-examples-2.2.0.jar : to get more information

Now suppose we want to run the wordcount example.

First, put the input data on HDFS (you have to create your own input.txt first)

hdfs dfs -put input.txt /input.txt

Next, execute the wordcount example

hadoop jar hadoop-mapreduce-examples-2.2.0.jar wordcount /input.txt /test_out

Finally, get the results

hadoop dfs -get /test_out test_out

The result file part-r-00000 show up in the directory test_out